

StrategicGrowthAnalysis-UniqueGiftsLtd

Downloading dataset

```
In [73]: import requests # for downloading the dataset
import os # for file operations
from tqdm import tqdm # progress bar library

url = "https://docs.google.com/spreadsheets/d/1RZ0nYqAUgSivbfIiMsbZfEy2yg5Kz"
file_path = "../data/dataset.csv" # Path to save the dataset

os.makedirs("../data", exist_ok=True)

response = requests.get(url, stream=True)
total_size = int(response.headers.get('content-length', 0))
block_size = 1024 # 1 KB

print("Downloading dataset with progress:")

with open(file_path, "wb") as file, tqdm(
    desc=file_path,
    total=total_size,
    unit='iB',
    unit_scale=True,
    unit_divisor=1024,
) as bar:
    for data in response.iter_content(block_size):
        file.write(data)
        bar.update(len(data))

if os.path.exists(file_path):
    print("Download complete!")
else:
    print("Download failed. Please check the URL or your internet connection")
```

Downloading dataset with progress:

../data/dataset.csv: 89.7MiB [01:13, 1.28MiB/s]

Download complete!

create a function to load the dataset

```
In [74]: import pandas as pd

if not os.path.exists(file_path): # Check if the file exists
    raise FileNotFoundError(f"Dataset file not found at {file_path}. Please

df = pd.read_csv(file_path) # Load the dataset
print(f"Dataset loaded successfully with {df.shape[0]} rows and {df.shape[1]}
```

Dataset loaded successfully with 1067371 rows and 8 columns.

your task is to analyze the dataset and provide insights on how UniqueGiftsLtd can strategically grow its business.

```
In [75]: # Display the first few rows of the dataset
try:
    df = pd.read_csv('../data/dataset.csv') # Adjust the path as necessary
    print("Dataset loaded successfully!")
except FileNotFoundError:
    print("Error: The file 'online_retail.csv' was not found in the 'data' f
    print("Please make sure you have downloaded the dataset and placed it in

# Display the first 5 rows to get an idea of the data
df.head()
```

Dataset loaded successfully!

```
Out[75]:
```

	Invoice	StockCode	Description	Quantity	InvoiceDate	Price	Customer ID
0	489434	85048	15CM CHRISTMAS GLASS BALL 20 LIGHTS	12	2009-12-01 07:45:00	6.95	13085.0
1	489434	79323P	PINK CHERRY LIGHTS	12	2009-12-01 07:45:00	6.75	13085.0
2	489434	79323W	WHITE CHERRY LIGHTS	12	2009-12-01 07:45:00	6.75	13085.0
3	489434	22041	RECORD FRAME 7" SINGLE SIZE	48	2009-12-01 07:45:00	2.10	13085.0
4	489434	21232	STRAWBERRY CERAMIC TRINKET BOX	24	2009-12-01 07:45:00	1.25	13085.0

Dataframe info

```
In [76]: print("Dataframe Info:")
df.info()

print("\n" + "="*50 + "\n")

# Get summary statistics for numerical columns
print("Numerical Describe:")
print(df.describe())
print("\n" + "="*50 + "\n")

# Count missing values in each column
print("Missing Values Count:")
print(df.isnull().sum())
```

Dataframe Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 1067371 entries, 0 to 1067370

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	Invoice	1067371 non-null	object
1	StockCode	1067371 non-null	object
2	Description	1062989 non-null	object
3	Quantity	1067371 non-null	int64
4	InvoiceDate	1067371 non-null	object
5	Price	1067371 non-null	float64
6	Customer ID	824364 non-null	float64
7	Country	1067371 non-null	object

dtypes: float64(2), int64(1), object(5)

memory usage: 65.1+ MB

Numerical Describe:

	Quantity	Price	Customer ID
count	1.067371e+06	1.067371e+06	824364.000000
mean	9.938898e+00	4.649388e+00	15324.638504
std	1.727058e+02	1.235531e+02	1697.464450
min	-8.099500e+04	-5.359436e+04	12346.000000
25%	1.000000e+00	1.250000e+00	13975.000000
50%	3.000000e+00	2.100000e+00	15255.000000
75%	1.000000e+01	4.150000e+00	16797.000000
max	8.099500e+04	3.897000e+04	18287.000000

Missing Values Count:

Invoice	0
StockCode	0
Description	4382
Quantity	0
InvoiceDate	0
Price	0
Customer ID	243007
Country	0

dtype: int64

Handle Duplicates & Missing Data)

```
In [77]: # Print shape before dropping duplicates
print(f"Shape before dropping duplicates: {df.shape}")

print(df.duplicated().sum())

# Drop duplicate rows
df.drop_duplicates(inplace=True) # re df

# Print shape after dropping duplicates
```

```

print(f"Shape after dropping duplicates: {df.shape}")

# Drop rows where Customer ID is missing
df.dropna(subset=['Customer ID'], inplace=True)

# Print shape after dropping missing Customer ID
print(f"Shape after dropping duplicates: {df.shape}")

# Verify that missing Customer IDs are handled
print("\nMissing values after handling Customer ID:")
print(df.isnull().sum())

```

Shape before dropping duplicates: (1067371, 8)
34335

Shape after dropping duplicates: (1033036, 8)

Shape after dropping duplicates: (797885, 8)

Missing values after handling Customer ID:

Invoice	0
StockCode	0
Description	0
Quantity	0
InvoiceDate	0
Price	0
Customer ID	0
Country	0

dtype: int64

Clean Transactional Data

```

In [78]: # Remove cancelled orders (Invoice starts with 'C')
df = df[~df['Invoice'].astype(str).str.startswith('C')]

# Ensure quantity is positive
df = df[df['Quantity'] > 0]

# Remove records where price is 0
df = df[df['Price'] > 0]

# Check the shape of the dataframe after cleaning
print(f"Shape after cleaning transactions: {df.shape}")

```

Shape after cleaning transactions: (779425, 8)

Filtering Non-Product Charges

```

In [79]: # (Assuming 'df' is your dataframe after the initial cleaning)

def is_product_code(code):
    """
    Checks if a stock code is likely a product by seeing if it contains any
    Returns True if it contains a digit, False otherwise.
    e.g., is_product_code('85123A') -> True
    """

```

```

        is_product_code('POST')    -> False
    """
    code_str = str(code)
    return any(char.isdigit() for char in code_str)
# =====

# --- Identify codes that would be removed using this new logic ---
all_unique_codes = df['StockCode'].unique()
codes_to_be_removed = [code for code in all_unique_codes if not is_product_code(code)]
print(f"Following codes (without any digits) will be removed: {codes_to_be_removed}")

# --- Filtering Step ---
print(f"\nShape before filtering: {df.shape}")

# Apply the function to the 'StockCode' column and keep only the rows that remain
df = df[df['StockCode'].apply(is_product_code)]

print(f"Shape after filtering: {df.shape}")

```

Following codes (without any digits) will be removed: ['POST', 'M', 'BANK CHARGES', 'PADS', 'ADJUST', 'D', 'DOT']

Shape before filtering: (779425, 8)

Shape after filtering: (776840, 8)

Create TotalPrice column

```

In [80]: # Create TotalPrice column
df['TotalPrice'] = df['Quantity'] * df['Price']

print(df.shape)
print(df.head())

```

(776840, 9)

	Invoice	StockCode	Description	Quantity	\
0	489434	85048	15CM CHRISTMAS GLASS BALL 20 LIGHTS	12	
1	489434	79323P	PINK CHERRY LIGHTS	12	
2	489434	79323W	WHITE CHERRY LIGHTS	12	
3	489434	22041	RECORD FRAME 7" SINGLE SIZE	48	
4	489434	21232	STRAWBERRY CERAMIC TRINKET BOX	24	

	InvoiceDate	Price	Customer ID	Country	TotalPrice
0	2009-12-01 07:45:00	6.95	13085.0	United Kingdom	83.4
1	2009-12-01 07:45:00	6.75	13085.0	United Kingdom	81.0
2	2009-12-01 07:45:00	6.75	13085.0	United Kingdom	81.0
3	2009-12-01 07:45:00	2.10	13085.0	United Kingdom	100.8
4	2009-12-01 07:45:00	1.25	13085.0	United Kingdom	30.0

Convert InvoiceDate to datetime

```
In [81]: # Convert InvoiceDate to datetime
df['InvoiceDate'] = pd.to_datetime(df['InvoiceDate'])

print(type(df['InvoiceDate']))
print(df['InvoiceDate'].dtype)

# 2. create new columns - Year, Month, DayOfWeek, HourOfDay
df['Year'] = df['InvoiceDate'].dt.year
df['Month'] = df['InvoiceDate'].dt.month
df['DayOfWeek'] = df['InvoiceDate'].dt.dayofweek      # Monday=0, Sunday=6
df['HourOfDay'] = df['InvoiceDate'].dt.hour

# 3. view the new columns (first 5 rows)
print(df[['InvoiceDate', 'Year', 'Month', 'DayOfWeek', 'HourOfDay']].head())
```

```
<class 'pandas.core.series.Series'>
datetime64[ns]
```

	InvoiceDate	Year	Month	DayOfWeek	HourOfDay
0	2009-12-01 07:45:00	2009	12	1	7
1	2009-12-01 07:45:00	2009	12	1	7
2	2009-12-01 07:45:00	2009	12	1	7
3	2009-12-01 07:45:00	2009	12	1	7
4	2009-12-01 07:45:00	2009	12	1	7

Convert Customer ID to integer

```
In [82]: df['Customer ID'] = df['Customer ID'].astype(int) # Convert Customer ID to
print(df['Customer ID'].dtype)
```

```
int64
```

Convert StockCode to string

```
In [83]: # Convert StockCode to string
df['StockCode'] = df['StockCode'].astype(str) # Convert StockCode to string
print(df['StockCode'].dtype)
```

```
object
```

create cleaned dataset - .csv file

```
In [84]: df.to_csv('../data/cleaned-dataset.csv', index=False)
```

Exploratory Data Analysis (EDA) & Insight Generation

```

In [85]: import matplotlib.pyplot as plt
import seaborn as sns

# Visualization style
sns.set_theme(style="whitegrid")
plt.rcParams['figure.figsize'] = (10, 6)

# Dataset load (Cleaned dataset)
import pandas as pd

# Robustly load and standardize the cleaned dataset (handles varying column
print("Loading cleaned dataset (robust)...")
df = pd.read_csv('../data/cleaned-dataset.csv', low_memory=False)

# --- Find and parse invoice date column ---
date_col = None
for c in df.columns:
    if 'invoice' in c.lower() and 'date' in c.lower():
        date_col = c
        break
if not date_col:
    # fallback common name
    if 'InvoiceDate' in df.columns:
        date_col = 'InvoiceDate'

if not date_col:
    raise ValueError(f"No invoice date column found in cleaned-dataset.csv.")

df[date_col] = pd.to_datetime(df[date_col], errors='coerce', dayfirst=True)
df.rename(columns={date_col: 'InvoiceDate'}, inplace=True)

# --- Standardize important column names (handles variants from different da
rename_map = {}
for c in df.columns:
    lc = c.lower().replace(' ', '').replace('_', '')
    if lc in ('invoiceno', 'invoice'):
        rename_map[c] = 'Invoice'
    if lc in ('invoicedate',):
        rename_map[c] = 'InvoiceDate'
    if lc in ('unitprice', 'price', 'unitpriceinc'):
        rename_map[c] = 'Price'
    if lc in ('customerid', 'customerid.'):
        rename_map[c] = 'Customer ID'
    if lc in ('stockcode',):
        rename_map[c] = 'StockCode'
    if lc in ('quantity',):
        rename_map[c] = 'Quantity'
    if lc in ('description',):
        rename_map[c] = 'Description'
    if lc in ('country',):
        rename_map[c] = 'Country'

df.rename(columns=rename_map, inplace=True)

# Verify required columns exist

```

```

required = ['Invoice', 'InvoiceDate', 'Price', 'Customer ID', 'Quantity']
missing = [r for r in required if r not in df.columns]
if missing:
    raise ValueError(f"Missing required columns in cleaned-dataset.csv: {missing}")

# Ensure numeric types
df['Quantity'] = pd.to_numeric(df['Quantity'], errors='coerce').fillna(0).astype(int)
df['Price'] = pd.to_numeric(df['Price'], errors='coerce').fillna(0.0)

# Ensure TotalPrice exists
if 'TotalPrice' not in df.columns:
    df['TotalPrice'] = df['Quantity'] * df['Price']

print(f"Dataset loaded and standardized. Shape: {df.shape}")

```

Loading cleaned dataset (robust)...

Dataset loaded and standardized. Shape: (776840, 13)

Temporal Analysis

```

In [86]: # Year-Month column
df['YearMonth'] = df['InvoiceDate'].dt.to_period('M')

# Group by Year-Month and sum TotalPrice
monthly_sales = df.groupby('YearMonth')['TotalPrice'].sum().reset_index()
monthly_sales['YearMonth'] = monthly_sales['YearMonth'].astype(str) # Convert to string

# Line chart Plot
plt.figure(figsize=(12, 6))
sns.lineplot(x='YearMonth', y='TotalPrice', data=monthly_sales, marker="o")
plt.title("Monthly Sales Trend")
plt.xlabel("Month")
plt.ylabel("Total Sales Revenue (£)")
plt.xticks(rotation=45)
plt.show()

```




Temporal Analysis – February & April Drops Check

```
In [87]: df['InvoiceDate'] = pd.to_datetime(df['InvoiceDate'])

# Extract Year and Month from InvoiceDate
df['Year'] = df['InvoiceDate'].dt.year
df['Month'] = df['InvoiceDate'].dt.month
df['MonthName'] = df['InvoiceDate'].dt.strftime('%B')

# Group by Year and MonthName to get total sales
monthly_sales_yearly = df.groupby(['Year', 'MonthName'])['TotalPrice'].sum()

# Reorder the months for better visualization
month_order = ['January', 'February', 'March', 'April', 'May', 'June',
               'July', 'August', 'September', 'October', 'November', 'December']
monthly_sales_yearly['MonthName'] = pd.Categorical(monthly_sales_yearly['MonthName'],
                                                    categories=month_order,
                                                    ordered=True)
monthly_sales_yearly = monthly_sales_yearly.sort_values(['Year', 'MonthName'])

# Plotting the monthly sales by year
plt.figure(figsize=(12,6))
sns.lineplot(data=monthly_sales_yearly, x='MonthName', y='TotalPrice', hue='Year')
plt.title("Monthly Sales by Year (Check for Feb & Apr Dips)")
plt.xlabel("Month")
plt.ylabel("Total Sales Revenue (£)")
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()

# February data
feb_sales = monthly_sales_yearly[monthly_sales_yearly['MonthName'] == "February"]
```

```

# April data
apr_sales = monthly_sales_yearly[monthly_sales_yearly['MonthName'] == "April"]

# Mean & std check
feb_mean = feb_sales.mean()
apr_mean = apr_sales.mean()
overall_mean = monthly_sales_yearly['TotalPrice'].mean()


print(f"February Mean Sales: {feb_mean}")
print(f"April Mean Sales: {apr_mean}")
print(f"Overall Mean Sales: {overall_mean}")

# Check pattern: True if all years' Feb sales < overall mean
if hasattr(feb_sales, 'empty') and feb_sales.empty:
    feb_pattern = False
else:
    feb_pattern = bool((feb_sales < overall_mean).all())

if hasattr(apr_sales, 'empty') and apr_sales.empty:
    apr_pattern = False
else:
    apr_pattern = bool((apr_sales < overall_mean).all())

print(f"February consistent dip? {feb_pattern}")
print(f"April consistent dip? {apr_pattern}")

```



February Mean Sales: 163228.31
 April Mean Sales: 197887.25333333333
 Overall Mean Sales: 203882.48628571432
 February consistent dip? False
 April consistent dip? False

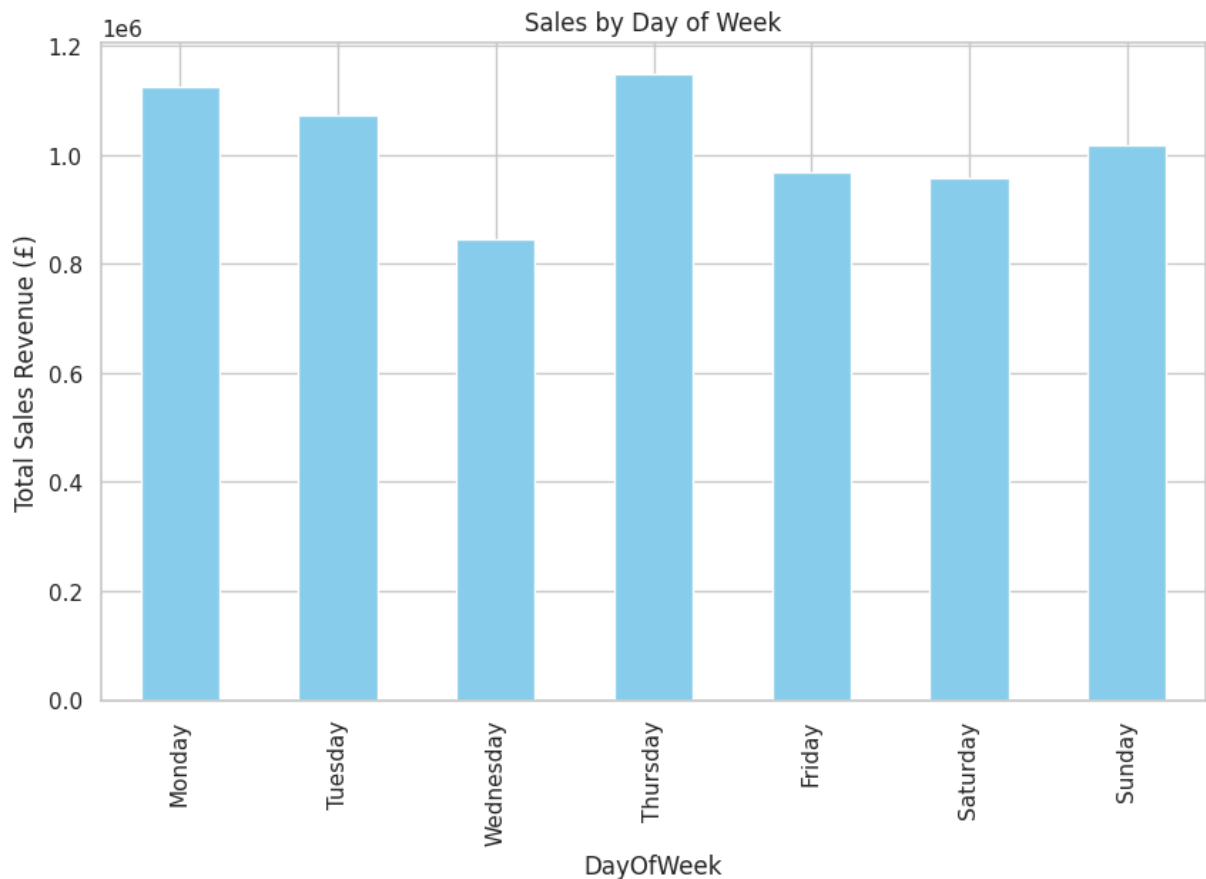
Temporal Analysis – Day & Hour Pattern


```
In [88]: df['DayOfWeek'] = df['InvoiceDate'].dt.day_name()
df['HourOfDay'] = df['InvoiceDate'].dt.hour

# Sales by Day
sales_by_day = df.groupby('DayOfWeek')['TotalPrice'].sum().reindex(
    ["Monday", "Tuesday", "Wednesday", "Thursday", "Friday", "Saturday", "Sunday"]
)

# Bar chart – Day
sales_by_day.plot(kind='bar', color='skyblue')
plt.title("Sales by Day of Week")
plt.ylabel("Total Sales Revenue (£)")
plt.show()

# Sales by Hour
sales_by_hour = df.groupby('HourOfDay')['TotalPrice'].sum()
sales_by_hour.plot(kind='bar', color='lightgreen')
plt.title("Sales by Hour of Day")
plt.ylabel("Total Sales Revenue (£)")
plt.show()
```





Geographic Analysis – Top 10 Countries

```
In [89]: country_sales = df.groupby('Country')['TotalPrice'].sum().sort_values(ascending=False)


# Top 10
top_10_countries = country_sales.head(10)

# Plot
sns.barplot(x='TotalPrice', y='Country', data=top_10_countries, palette="viridis")
plt.title("Top 10 Countries by Sales Revenue")
plt.xlabel("Total Sales Revenue (£)")
plt.ylabel("Country")
plt.show()
```

/tmp/ipykernel_5583/162923314.py:7: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='TotalPrice', y='Country', data=top_10_countries, palette="viridis")
```



Geographic Analysis – UK vs Other Countries

```
In [90]: uk_revenue = df[df['Country'] == 'United Kingdom']['TotalPrice'].sum()
total_revenue = df['TotalPrice'].sum()
uk_percentage = (uk_revenue / total_revenue) * 100
print(f"UK Revenue %: {uk_percentage:.2f}%")
```



UK Revenue %: 83.67%

DayOfWeek and HourOfDay bar charts

```
In [91]: import matplotlib.pyplot as plt

# Revenue by DayOfWeek
day_sales = df.groupby('DayOfWeek')['TotalPrice'].sum()
day_sales.plot(kind='bar', figsize=(7,4), title="Revenue by Day of Week")
plt.ylabel("Revenue")
plt.show()

# Revenue by HourOfDay
hour_sales = df.groupby('HourOfDay')['TotalPrice'].sum()
hour_sales.plot(kind='bar', figsize=(7,4), title="Revenue by Hour of Day")
plt.ylabel("Revenue")
plt.show()
```



Product Performance Analysis

```
In [92]: # Top 10 by Quantity Sold
top_quantity = df.groupby('Description')['Quantity'].sum().sort_values(ascending=False)

# Top 10 by Revenue
top_revenue = df.groupby('Description')['TotalPrice'].sum().sort_values(ascending=False)

print("Top 10 by Quantity Sold:\n", top_quantity)
print("\nTop 10 by Revenue:\n", top_revenue)
```

Top 10 by Quantity Sold:

Description	Quantity
WORLD WAR 2 GLIDERS ASSTD DESIGNS	105185
WHITE HANGING HEART T-LIGHT HOLDER	91757
PAPER CRAFT , LITTLE BIRDIE	80995
ASSORTED COLOUR BIRD ORNAMENT	78234
MEDIUM CERAMIC TOP STORAGE JAR	77916
JUMBO BAG RED RETROSPOT	74224
BROCADE RING PURSE	70082
PACK OF 60 PINK PAISLEY CAKE CASES	54592
60 TEATIME FAIRY CAKE CASES	52828
PACK OF 72 RETRO SPOT CAKE CASES	45129

Name: Quantity, dtype: int64

Top 10 by Revenue:

Description	TotalPrice
REGENCY CAKESTAND 3 TIER	277656.25
WHITE HANGING HEART T-LIGHT HOLDER	247048.01
PAPER CRAFT , LITTLE BIRDIE	168469.60
JUMBO BAG RED RETROSPOT	134307.44
ASSORTED COLOUR BIRD ORNAMENT	124351.86
PARTY BUNTING	103283.38
MEDIUM CERAMIC TOP STORAGE JAR	81416.73
PAPER CHAIN KIT 50'S CHRISTMAS	76598.18
CHILLI LIGHTS	69084.30
JUMBO BAG STRAWBERRY	64127.77

Name: TotalPrice, dtype: float64

3.3 Phase 3: Advanced Analytics - RFM Customer Segmentation

```
In [93]: import pandas as pd
import numpy as np

# Use the already-loaded `df` (from earlier EDA). If not present, load robustly
try:
    df
except NameError:
    df = pd.read_csv('../data/cleaned-dataset.csv', low_memory=False)

# Ensure InvoiceDate is datetime
if df['InvoiceDate'].dtype == object or not np.issubdtype(df['InvoiceDate'], np.datetime64):
    df['InvoiceDate'] = pd.to_datetime(df['InvoiceDate'], errors='coerce', utc=True)
```

```

# Ensure TotalPrice exists
if 'TotalPrice' not in df.columns:
    df['TotalPrice'] = pd.to_numeric(df['Quantity'], errors='coerce').fillna(0)

# Snapshot date = 1 day after last InvoiceDate
snapshot_date = df['InvoiceDate'].max() + pd.Timedelta(days=1)
print(f"Snapshot date: {snapshot_date}")

# Recency
recency_df = df.groupby('Customer ID').agg(Recency=('InvoiceDate', lambda x:

# Frequency (number of unique invoices)
frequency_df = df.groupby('Customer ID').agg(Frequency=('Invoice', 'nunique')

# Monetary (total spent)
monetary_df = df.groupby('Customer ID').agg(Monetary=('TotalPrice', 'sum')).

# Merge
rfm_df = recency_df.merge(frequency_df, on='Customer ID').merge(monetary_df,

# Add average order value and average order quantity to help classify wholes
# avg_order_value = Monetary / Frequency
rfm_df['AvgOrderValue'] = rfm_df['Monetary'] / rfm_df['Frequency']

# average quantity per invoice per customer
invoice_qty = df.groupby(['Customer ID', 'Invoice']).agg(InvoiceQty=('Quantity', 'sum'))
avg_qty = invoice_qty.groupby('Customer ID').agg(AvgOrderQty=('InvoiceQty', 'mean'))
rfm_df = rfm_df.merge(avg_qty, on='Customer ID', how='left')

# RFM scoring using quintiles. Use rank for frequency to reduce qcut issues
# Handle small number of customers by using qcut with duplicates dropped when
r_labels = [5,4,3,2,1]
f_labels = [1,2,3,4,5]
m_labels = [1,2,3,4,5]

try:
    rfm_df['R_Score'] = pd.qcut(rfm_df['Recency'], 5, labels=r_labels)
    rfm_df['F_Score'] = pd.qcut(rfm_df['Frequency'].rank(method='first'), 5,
                                labels=f_labels)
    rfm_df['M_Score'] = pd.qcut(rfm_df['Monetary'], 5, labels=m_labels)
except ValueError:
    # fallback when there are too few unique values
    rfm_df['R_Score'] = pd.cut(rfm_df['Recency'], bins=5, labels=r_labels)
    rfm_df['F_Score'] = pd.cut(rfm_df['Frequency'].rank(method='first'), bins=5,
                                labels=f_labels)
    rfm_df['M_Score'] = pd.cut(rfm_df['Monetary'], bins=5, labels=m_labels)

# Convert to int (coerce if NaN)
rfm_df['R_Score'] = pd.to_numeric(rfm_df['R_Score'], errors='coerce').fillna(0)
rfm_df['F_Score'] = pd.to_numeric(rfm_df['F_Score'], errors='coerce').fillna(0)
rfm_df['M_Score'] = pd.to_numeric(rfm_df['M_Score'], errors='coerce').fillna(0)

# RFM segment code and description
rfm_df['RFM_Segment'] = rfm_df['R_Score'].astype(str) + rfm_df['F_Score'].astype(str) + rfm_df['M_Score'].astype(str)

def map_rfm_segment(row):
    r, f, m = row['R_Score'], row['F_Score'], row['M_Score']
    if r >= 4 and f >= 4 and m >= 4:

```



```

        return 'Champions'
    if f >= 4 and r >= 3:
        return 'Loyal Customers'
    if r >= 4 and f <= 2:
        return 'Recent Low-Frequency'
    if r <= 2 and f <= 2:
        return 'Hibernating'
    if r <= 2 and f >= 3:
        return 'At-Risk Customers'
    return 'Potential Loyalists'

rfm_df['RFM_Description'] = rfm_df.apply(map_rfm_segment, axis=1)

# Classify Wholesaler vs Retail based on simple heuristics
# Thresholds can be tuned; these are initial suggestions per project spec
wholesale_qty_threshold = 20 # avg qty per order
wholesale_value_threshold = 200.0 # avg order value

rfm_df['CustomerType'] = np.where((rfm_df['AvgOrderQty'] > wholesale_qty_thr

# Summary comparisons
summary = rfm_df.groupby('CustomerType').agg(
    Customers=('Customer ID', 'unique'),
    AvgFrequency=('Frequency', 'mean'),
    AvgMonetary=('Monetary', 'mean'),
    AvgOrderQty=('AvgOrderQty', 'mean'),
    AvgOrderValue=('AvgOrderValue', 'mean')
).reset_index()

print('\nCustomer type summary:')
print(summary)

# Save RFM segments
rfm_df.to_csv('../data/rfm_segments.csv', index=False)
print('\nSaved RFM segments to ../data/rfm_segments.csv')

```

Snapshot date: 2011-12-11 17:19:00

Customer type summary:

	CustomerType	Customers	AvgFrequency	AvgMonetary	AvgOrderQty	\
0	Retail	125	2.720000	140.627840	11.399038	
1	Wholesaler	5728	6.332577	2979.138843	256.747940	

	AvgOrderValue
0	61.894360
1	388.211929

Saved RFM segments to ../data/rfm_segments.csv

Map RFM scores to descriptive segments

```

In [94]: # Define segment mapping
def rfm_segment(row):

```

```

r, f, m = row['R_Score'], row['F_Score'], row['M_Score']
if r >= 4 and f >= 4 and m >= 4:
    return 'Champions'
elif r >= 4 and f >= 4:
    return 'Loyal Customers'
elif r >= 4 and f <= 2:
    return 'New Customers'
elif r <= 2 and f >= 3:
    return 'At-Risk Customers'
else:
    return 'Hibernating'

rfm_df['Segment'] = rfm_df.apply(rfm_segment, axis=1)
rfm_df['Segment'].value_counts()

```

```

Out[94]: Segment
Hibernating      3024
At-Risk Customers 1376
Champions        1093
New Customers     227
Loyal Customers   133
Name: count, dtype: int64

```

3.4 Data Enrichment (API) - example: get country ISO codes using REST Countries API

```

In [95]: import requests
import time

countries = df['Country'].dropna().unique().tolist()
country_rows = []
for country in countries:
    try:
        resp = requests.get(f'https://restcountries.com/v3.1/name/{country}?')
        if resp.status_code == 200:
            data = resp.json()
            if isinstance(data, list) and len(data) > 0:
                c = data[0]
                iso2 = c.get('cca2')
                iso3 = c.get('cca3')
                country_rows.append({'Country': country, 'ISO2': iso2, 'ISO3': iso3})
            else:
                country_rows.append({'Country': country, 'ISO2': None, 'ISO3': None})
        except Exception:
            country_rows.append({'Country': country, 'ISO2': None, 'ISO3': None})
    time.sleep(0.15)

country_codes = pd.DataFrame(country_rows)
country_codes.to_csv('../data/country_codes_enrichment.csv', index=False)
print('\nSaved country enrichment to ../data/country_codes_enrichment.csv')

```


```
# Merge example (non-destructive):  
# df = df.merge(country_codes, on='Country', how='left')
```

Saved country enrichment to ../data/country_codes_enrichment.csv

Investigate wholesaler hypothesis

```
In [96]: # Distribution of Monetary value per customer  
rfm_df['Monetary'].plot(kind='hist', bins=50, figsize=(8,5), title="Customer  
plt.xlabel("Monetary Value")  
plt.show()  
  
rfm_df['Monetary'].plot(kind='box', figsize=(6,4), title="Monetary Value per  
plt.show()
```





Currency conversion with API + add USD/EUR columns

```
In [100... import requests

# Use a free API that doesn't require authentication
url = "https://api.fxratesapi.com/latest?base=GBP&currencies=USD,EUR"

try:
    response = requests.get(url)
    response.raise_for_status()
    data = response.json()

    if 'rates' in data:
        gbp_to_usd = data['rates']['USD']
        gbp_to_eur = data['rates']['EUR']
        print(f"Live rates - GBP to USD: {gbp_to_usd}, GBP to EUR: {gbp_to_eur}")
    else:
        # Fallback rates
        gbp_to_usd = 1.27
        gbp_to_eur = 1.17
        print("Using fallback rates")

except requests.RequestException as e:
    print(f"API request failed: {e}")
    # Fallback rates
    gbp_to_usd = 1.27
    gbp_to_eur = 1.17
    print("Using fallback rates")
```

```

# Apply currency conversion
df['TotalPrice_USD'] = df['TotalPrice'] * gbp_to_usd
df['TotalPrice_EUR'] = df['TotalPrice'] * gbp_to_eur

# Display results
print(df[['Invoice', 'TotalPrice', 'TotalPrice_USD', 'TotalPrice_EUR']].head)

```

Live rates - GBP to USD: 1.34921797, GBP to EUR: 1.153622

	Invoice	TotalPrice	TotalPrice_USD	TotalPrice_EUR
0	489434	83.40	112.524779	96.212075
1	489434	81.00	109.286656	93.443382
2	489434	81.00	109.286656	93.443382
3	489434	100.80	136.001171	116.285098
4	489434	30.00	40.476539	34.608660
..
95	489442	13.50	18.214443	15.573897
96	489442	14.85	20.035887	17.131287
97	489442	19.80	26.714516	22.841716
98	489442	15.00	20.238270	17.304330
99	489442	17.70	23.881158	20.419109

[100 rows x 4 columns]